

Copilot in the Classroom

Adopt Fast, Govern Faster

Why institutions need a repeatable way to decide, govern, and answer for AI in education.

Ahlam Ansari

Research Scholar, KJ Somaiya School of Engineering

Assistant Professor, MH Saboo Siddik College of Engineering

The pitch is winning

Every vendor conversation lands on the same three promises.

1

Personalised Learning

Adaptive paths for every student, on paper.

2

Time Savings

Hours back for teachers buried in admin work.

3

Closed Achievement Gaps

The metric every board wants to see move.

Institutions say yes. Almost none have a repeatable process for deciding whether they should — or a plan for the day the tool gets something wrong.

THE STORY Picture the pitch deck: an AI coding-tutor platform promising each student a ‘personalised path’ through data structures and algorithms, a 30% cut in the hours faculty spend grading code, and a projected close on the programming pass-rate gap within two years. No pilot data yet — just projections. That’s the deck that gets approved in a 45-minute curriculum committee meeting.



I've watched that gap hurt institutions more than once.

THE STORY A university I worked with rolled out an AI code-review and auto-grading tool mid-semester in a Data Structures course. Three weeks in, a TA noticed it was scoring international students' submissions lower — not for logic or correctness, but for variable-naming and comment style that didn't match a 'standard' convention. Nobody had asked, before launch, who was supposed to catch that.

The pattern is boring and consistent

The tool ships before anyone in the room asks:

1

Who audits the outputs?

Once it's live, who is actually checking what it decides — and how often?

STORY *One university's auto-grading platform ran unaudited for eighteen months across several programming courses — a review then found a 12% false-positive rate on plagiarism flags for independently written code.*

2

What happens to student data?

Where it goes, who can see it, and how long it's kept after the pilot ends.

STORY *A popular coding-practice platform's contract let it retain students' submitted code and IDE activity logs for 'product improvement,' no deletion clause — most faculty didn't know until a student asked.*

3

What do we say when we can't explain it?

The call that lands on a student, a department head, or an integrity committee.

STORY *A plagiarism detector flagged a student's assignment as 90% similar to another's — based only on shared starter code every student was given — and couldn't give a real-time, plain-language reason why in the hearing.*

What does the department say to a student — or a committee, or a regulator — when the AI makes a call no one in the room can actually explain?

THIS IS THE QUESTION THAT GOVERNANCE ANSWERS — NOT THE PITCH DECK

THE STORY *Imagine an AI proctoring tool on an online coding exam flags a student for a ‘suspicious keystroke pattern’ — they paused to think, or used an IDE shortcut. The student sits before an academic-integrity committee asking why. The honest answer — ‘the model flagged an unusual typing rhythm and we don’t fully know why’ — is not one any faculty member wants to give in that room.*

An evaluation framework built for schools, colleges & universities

Three pillars, built specifically for K–12 and higher ed — not repurposed enterprise checklists.

1

Transparency

Can the vendor show, in plain language, how the system reaches a decision?

STORY

Ask the vendor: ‘Walk me through why Student A’s submission got flagged for plagiarism and Student B’s didn’t.’ If the answer is ‘the model just knows,’ that’s your answer too.

2

Bias Risk

Has it been tested against your student population — not a generic benchmark?

STORY

An online coding-exam proctoring tool tuned for stable, well-lit conditions flagged students on shared campus Wi-Fi or older webcams for ‘suspicious activity’ at far higher rates — surfaced only after deployment.

3

Data Accountability

Who owns the data, where does it live, and what happens when the contract ends?

STORY

Ask what happens to a student’s submitted source code and project repos the day they graduate or transfer — most coding-platform contracts are silent on exactly that.

Every pillar maps to one question: could you explain this decision to someone outside the room?

Bring them in before, not after

A stakeholder plan that puts every voice at the table before the decision is made.

1

Faculty

What changes in grading and labs, and where their judgment still overrides the tool.

STORY

Give faculty and TAs a one-page 'override log' — a record of when they overruled the auto-grader's or plagiarism-detector's flag and why. It protects them and builds your bias-audit trail for free.

2

Administrators

What they're accountable for, and what they need on hand before a department asks.

STORY

Before rollout, prep a one-paragraph answer to 'why did we choose this auto-grading vendor' a department head could repeat verbatim. If you can't write it, you're not ready to launch.

3

Students

What data is collected, why, and how to raise a concern — in plain language.

STORY

Send a plain-language notice before the semester starts: what tool grades or checks their code, what data it collects — including IDE activity — and a named contact for concerns — not a clause buried in the syllabus.

Same tool, three different conversations — each one before the rollout, not after a complaint.

The framework in practice

Running the tools your CS students already use in the lab through the same three pillars.

1

GitHub Copilot

Business / Enterprise tier

Transparency **Strong**

Published policy docs, admin audit logs, org-wide model governance controls.

Bias Risk **Mixed**

Fairness is research and published; but the experiments on bias detection is not.

Data **Strong***

Contractually exempt from AI-training use — *but free/individual tiers opted into training by default since Apr 2026.

2

Amazon Q Developer

CodeWhisperer lineage

Transparency **Mixed**

Enterprise configuration is documented; individual decision logic is not.

Bias Risk **Weak**

No independent, published bias audit for suggestion behavior.

Data **Mixed**

Opt-out by default for individuals; enterprises must explicitly configure full data controls.

3

Cursor / JetBrains AI

IDE-embedded assistants

Transparency **Mixed**

Markets 'privacy-first' positioning; underlying model logic stays with the API provider.

Bias Risk **Weak**

No independent bias audits published for either tool.

Data **Mixed**

Cursor routes through third-party APIs (e.g. OpenAI) — a second policy to track, not one.

Even the strongest score is a paid-tier promise, not a guarantee. Copilot's own 2026 shift on free-tier data is the case study for why governance is a habit, not a one-time checklist.


Strategy without a governance layer isn't leadership. It's exposure dressed up as progress.

Adopt fast. Govern faster. Let's build the layer that lets you answer for the tool — before you're asked to.

THE STORY *Callback to that integrity-committee hearing: the institutions that survive it aren't the ones with the fanciest AI — they're the ones who wrote the governance layer before they needed it.*

Contact Information

Questions and discussion welcome.



<https://www.linkedin.com/in/ahlamshakeel/>



@ahlamshakeel

Thank You

“

Algorithms are opinions embedded in code.

— Cathy O’Neil, Weapons of Math Destruction

Ahlam Ansari

Assistant Professor, Computer Engineering — Doctoral Researcher, Trustworthy AI

@ahlamshakeel | #DevDays #DevDaysMumbai #mumtechup